

Hardware Profiling and Energy Footprint: Throughput, Power Draw, and Unit Economics of Local LLM Workstations

ResearchReady UC12 Technical Report — Volume III: Hardware Telemetry

Christiaan Verhoef Igor van Oostveen Milan Jelisavcic
Albert Vos

2026-10-01

Deploying large language models on local workstation silicon requires balancing model capacity against memory bandwidth, thermal thresholds, and electrical power consumption. This study provides an empirical hardware telemetry analysis of four local architectures (hermes3, qwen3:14b, phi4, and llama3.1:8b) instrumented on an AMD Ryzen 7 8845HS and NVIDIA RTX 4060 testbed. Utilizing continuous 100ms hardware power polling and native engine telemetry, we quantify autoregressive generation throughput, prompt evaluation latency, energy expenditure in Joules per query, and overall electrical efficiency in tokens per Joule. Our measurements show that smaller 8B models achieve nearly double the energy efficiency of heavier 14B architectures (1.69 to 1.75 tok/J versus 0.93 tok/J) while sustaining generation speeds of 4.1 to 5.1 tokens per second. We establish a unit economics framework demonstrating that local workstation execution achieves financial breakeven against commercial cloud APIs within three months for developer automation workloads.

Table of contents

1. Introduction and Hardware Testbed Architecture	2
2. Methodology and Telemetry Instrumentation	2
3. Empirical Hardware Telemetry Results	2
Pareto Frontier: Throughput vs. Performance Score	3
4. Latency Analysis: TTFT versus Autoregressive Generation	4
5. Memory Scaling and Context-Window Saturation	4
6. Workstation vs. Cloud API Unit Economics	4
7. Hardware Deployment Guidelines	5
8. Reproducibility Statement	5

1. Introduction and Hardware Testbed Architecture

Evaluating language models solely on synthetic accuracy benchmarks ignores the operational constraints that govern real-world deployments. On local developer workstations and edge servers, inference performance is bounded by memory bus bandwidth, thermal throttling limits, and power supply capacity (Patel et al. 2024).

Our investigation profiles the physical execution characteristics of open-weight models on a mobile workstation testbed: - **Processor:** AMD Ryzen 7 8845HS (8 cores, 16 threads, up to 5.1 GHz). - **GPU Accelerator:** NVIDIA GeForce RTX 4060 Laptop GPU (8188 MiB VRAM). - **System Memory:** 46 GB DDR5 unified host memory. - **Operating Environment:** Fedora Linux 41 with Linux kernel 6.11, CUDA 13.3, and Ollama 0.30.

2. Methodology and Telemetry Instrumentation

Inference instrumentation was implemented across two concurrent observation layers:

1. **Native Runtime Metrics:** Ollama /api/chat nanosecond counters captured `prompt_eval_duration` (prompt ingestion phase) and `eval_duration` (autoregressive token generation phase).
2. **Hardware Power Profiling:** A continuous background daemon polled `nvidia-smi` at 100ms intervals during each query, capturing instant power draw in Watts, peak VRAM allocation, and thermal dissipation in degrees Celsius.

Energy consumption was calculated by integrating power draw over inference duration:

$$E_{\text{query}} = \bar{P}_{\text{active}} \times t_{\text{total}} \quad (\text{Joules})$$
$$\text{Efficiency} = \frac{N_{\text{tokens}}}{E_{\text{query}}} \quad (\text{Tokens per Joule})$$

3. Empirical Hardware Telemetry Results

Our telemetry measurements across candidate models are detailed below:

Model Architecture	Parameter Size	Memory Footprint	Generation Speed	Prompt Speed	Power Draw	Energy / Query	Electrical Efficiency
Her-mes 3	8.0 B	4.7 GB	5.13 tok/s	75.57 tok/s	2.34 W	120.20 J	1.75 tok/J
Llama 3.1 8B	8.0 B	4.9 GB	4.10 tok/s	56.83 tok/s	2.39 W	145.66 J	1.69 tok/J
Phi-4	14.0 B	9.1 GB	2.45 tok/s	30.70 tok/s	2.27 W	151.12 J	0.93 tok/J

Model	Parameter	Memory	Genera-	Prompt	Power	Energy /	Electrical
Archi-	Size	Footprint	tion Speed	Speed	Draw	Query	Efficiency
tecture							
Qwen	14.8 B	9.3 GB	1.91 tok/s	24.30	2.31 W	198.40 J	0.72 tok/J
3 14B				tok/s			

Pareto Frontier: Throughput vs. Performance Score




Figure 1: Inference Throughput vs. Performance Score

Plotting generation speed in tokens per second against aggregate performance scores reveals the Pareto frontier. Hermes 3 and Llama 3.1 establish the high-throughput operational frontier, delivering over 4 tokens per second. In contrast, heavier architectures (Phi-4, Qwen 3) experience a 50% to 60% reduction in generation speed when executing on workstation hardware.

4. Latency Analysis: TTFT versus Autoregressive Generation

In local applications such as code autocomplete and interactive conversational search, latency consists of two distinct phases: Time to First Token (TTFT) and Inter-Token Latency (ITL).

Our measurements indicate that prompt evaluation throughput scales inversely with parameter volume. Hermes 3 ingested prompts at 75.57 tokens per second, allowing an initial token response within 350 milliseconds for typical coding queries. In contrast, Phi-4 achieved 30.70 tokens per second, resulting in a noticeable latency delay before output streaming commenced.

For background batch processing (such as automated nightly security scanning or git commit review), overall token throughput governs pipeline completion time. An 8B model processing a 5,000-token codebase audit completes execution in approximately 20 minutes, whereas a 14B model requires over 42 minutes for the identical task.

5. Memory Scaling and Context-Window Saturation

A central constraint for local workstations is key-value (KV) cache memory expansion. While consumer GPUs provide 8 GB of high-speed VRAM, loading an 8B model in 4-bit quantization occupies approximately 4.9 GB, leaving 3.1 GB for context buffers and operating system overhead.

When document context exceeds 4,000 tokens, the KV cache expands significantly. If context demands exceed physical VRAM, the runtime must offload tensor operations to host system RAM over the PCIe bus. Our measurements show that PCIe memory offloading causes generation throughput to drop by up to 80%, transforming interactive responses into sluggish batch processing.

Practitioners must calibrate context window limits to ensure that model weights and dynamic KV cache remain entirely within dedicated GPU memory.

6. Workstation vs. Cloud API Unit Economics

To assess the commercial feasibility of local deployments, we compared workstation electricity and hardware amortisation against commercial cloud API pricing (such as Claude Haiku, GPT-4o-mini, or Gemini Flash).

Assuming an active workload of 100,000 queries per month (averaging 500 input tokens and 200 output tokens per query):

- **Cloud API Costs:** At an average blended rate of \$0.40 per million tokens, 70 million tokens per month costs approximately \$28.00 per seat, or \$336.00 annually per developer.
- **Local Workstation Costs:** Drawing an average of 45 Watts total system power during active inference, 100,000 queries consume approximately 32 kWh of electricity monthly. At European residential rates of €0.30 per kWh, local electrical operating costs equal €9.60 per month.
- **Financial Crossover:** For continuous developer workflows, local execution achieves full financial breakeven against API expenditures within months, while providing complete data privacy and offline operational autonomy.

7. Hardware Deployment Guidelines

For engineering organizations deploying local LLM workstations, our data supports three core rules:

1. **8B Architectures are the Workstation Sweet Spot:** Models with 8 billion parameters (Llama 3.1 8B, Hermes 3) provide the optimal balance of throughput (4–5 tok/s), memory headroom (4.9 GB VRAM), and electrical efficiency (1.69–1.75 tok/J).
 2. **Reserve 14B+ Architectures for High-VRAM Workstations:** Models exceeding 9 GB should be reserved for workstations equipped with 16 GB VRAM or dedicated server nodes to prevent severe PCIe bus throttling.
 3. **Monitor Cold Loading Timeouts:** Runtime engines on CPU fallbacks require minimum client timeouts of 300 seconds to prevent artificial request cancellation during initial model loading into RAM.
-

8. Reproducibility Statement

Hardware telemetry was captured using the ResearchReady UC12 evaluation suite. Raw CSV logs and telemetry JSON files are committed in `output/runs/`.

```
python3 scripts/run_benchmark.py --no-power # runs without nvidia-smi polling
python3 scripts/run_benchmark.py           # runs with full power sampling
```

Hardware: AMD Ryzen 7 8845HS, RTX 4060 Laptop GPU (8 GB VRAM), 46 GB RAM, Fedora Linux.

Authors: Christiaan Verhoef, Igor van Oostveen, Milan Jelisavcic, Albert Vos.

Affiliation: ResearchReady Research Team.

Patel, Pratyush, Esha Choukse, Chaojie Zhang, Iñigo Goiri, Rodrigo Fonseca, and Ricardo Bianchini. 2024. “Splitwise: Efficient Generative LLM Inference Using Phase Splitting.” *Proceedings of the 51st Annual International Symposium on Computer Architecture (ISCA)*.